

Text Clustering Example

- This example shows text Clustering as an example of unsupervised machine learning task.
- The main steps are:
 - Explore dataset
 - Data Preparation
 - Feature Engineering
 - Model Training
 - Performance Assessment
- Clustering: grouping similar objects together based on their inherent attributes.
- Being unsupervised approach there is no need of prelabelled datasets.
- Examples: KMeans, PAM, DBSCAN, Spectral clustering, etc.
- The choice of the algorithm mainly depends on the problem domain and the algorithm property. For example, to use KMeans we should have pre-knowledge of the number of clusters to obtain. However, if we do not know the number of clusters already, then DBSCAN is better.

Problem Definition

- As discussed above, the problem in context is unsupervised learning problem.
- We will use BCC news datasets available from <http://mlg.ucd.ie/files/datasets/bbc-fulltext.zip>. A total of 2225 documents with five different news categories(i.e. business, entertainment, politics, sport, and tech) are available in the dataset.

```
In [1]: # Import necessary libraries

import numpy as np
import pandas as pd
from sklearn.cluster import MiniBatchKMeans
from sklearn.cluster import KMeans
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
```

1. Load and explore data

- The dataset contains 5 different categories in 5 different folders.

```
In [2]: from sklearn.datasets import load_files

# for reproducibility
random_state = 22

DATA_DIR = "./data/bbc/"
data = load_files(DATA_DIR, encoding="utf-8", decode_error="replace", random_state=random_state)
df = pd.DataFrame(list(zip(data['data'], data['target'])), columns=['text', 'label'])
```

It is necessary to do an exploratory data analysis in order to gain some insights from the data.

How is the data like?

```
In [3]: df.head()
```

```
Out[3]:
```

	text	label
0	Jones files Conte lawsuit\n\nMarion Jones has ...	3
1	English clubs make Euro history\n\nAll four of...	3
2	Tarantino 'to make Friday sequel'\n\nDirector ...	1
3	Lesotho textile workers lose jobs\n\nSix forei...	0
4	Celtic make late bid for Bellamy\n\nNewcastle ...	3

how many 'class label' present in the dataset??

```
In [4]: df['label'].unique()
```


```
Out[4]: array([3, 1, 0, 4, 2])
```

Is the dataset balanced??

If the dataset is balanced then, the dataset contains an approximately equal portion of each class.

```
In [5]: df['label'].value_counts().plot(kind='bar')
```

```
Out[5]: <AxesSubplot:>
```



2. Feature Engineering

- we will use TFIDF approach.
- TfidfVectorizer will convert all the documents to a matrix of TF-IDF features.
- With Tfidfvectorizer we will calculate the word counts, IDF and TFIDF values in a go.

```
In [6]: tfidf_vectorizer = TfidfVectorizer(stop_words="english")  
tfidf_vectorizer.fit(df.text.values)  
features = tfidf_vectorizer.transform(df.text.values)
```

Lets check some values from the first document

```
In [7]: # get the first vector out (for the first document)  
first_vector_tfidfvectorizer=features[0]  
  
# place tf-idf values in a pandas data frame  
feature_df = pd.DataFrame(first_vector_tfidfvectorizer.T.todense(), index=tfidf_vectorizer.get_feature_names()  
feature_df.sort_values(by=["tfidf"],ascending=False).head()
```

```
Out[7]:
```

tfidf

	tfidf
conte	0.673159
jones	0.292039
marion	0.213606
doping	0.181978

3. Model Building

- We will check a version of KMeans algorithm, i.e. Mini-batch KMeans (MB KMeans).
- MB KMeans is similar to the standard KMeans. But in the MB KMeans the computationally expensive step is performed on a random sample only instead of the entire data.
- We will set the number of clusters = 5 because our dataset contains categories.
- But, if we do not have knowledge of the different labels then we may have to using domain knowledge and other ways to get the optimal numbers of clusters.

```
In [8]: mkmeans = MiniBatchKMeans(n_clusters=5, random_state=random_state)
mkmeans.fit(features)
```

```
Out[8]: MiniBatchKMeans(n_clusters=5, random_state=22)
```

lets check the predictions using labels_ attribute of the model.

```
In [9]: mkmeans.predict(features)
print (mkmeans.labels_)
print ("\n")
print(mkmeans.cluster_centers_)
```

```
[4 2 4 ... 4 4 1]
```

```
[[1.34473036e-04 1.09906331e-02 0.00000000e+00 ... 0.00000000e+00
 0.00000000e+00 0.00000000e+00]
 [1.68041680e-04 1.05754629e-02 0.00000000e+00 ... 0.00000000e+00
 0.00000000e+00 0.00000000e+00]
 [2.47895846e-04 2.12690116e-03 0.00000000e+00 ... 0.00000000e+00
```

```

4.60818743e-04 0.00000000e+00]
[0.00000000e+00 5.00195763e-03 0.00000000e+00 ... 0.00000000e+00
0.00000000e+00 0.00000000e+00]
[4.67908112e-06 1.15377847e-02 4.45013653e-05 ... 4.09596975e-05
0.00000000e+00 4.63992050e-04]]

```

Check KMeans

```

In [10]: kmeans = KMeans(n_clusters=5, init='k-means++', max_iter=100, n_init=1)
kmeans.fit(features)
kmeans.predict(features)
print (mkmeans.labels_)
print ("\n")
print(mkmeans.cluster_centers_)

[4 2 4 ... 4 4 1]

```

```

[[1.34473036e-04 1.09906331e-02 0.00000000e+00 ... 0.00000000e+00
0.00000000e+00 0.00000000e+00]
[1.68041680e-04 1.05754629e-02 0.00000000e+00 ... 0.00000000e+00
0.00000000e+00 0.00000000e+00]
[2.47895846e-04 2.12690116e-03 0.00000000e+00 ... 0.00000000e+00
4.60818743e-04 0.00000000e+00]
[0.00000000e+00 5.00195763e-03 0.00000000e+00 ... 0.00000000e+00
0.00000000e+00 0.00000000e+00]
[4.67908112e-06 1.15377847e-02 4.45013653e-05 ... 4.09596975e-05
0.00000000e+00 4.63992050e-04]]

```

```

In [11]: # lets order the centroids

order_centroids = kmeans.cluster_centers_.argsort()[:, :-1]
order_centroids

```

```

Out[11]: array([[29039, 5711, 27997, ..., 18590, 18591, 0],
[22845, 12210, 16592, ..., 16826, 16827, 0],
[17704, 22845, 15294, ..., 15629, 15628, 14562],
[11464, 9609, 22845, ..., 16998, 16999, 14562],
[10650, 3844, 3182, ..., 13381, 13382, 14562]])

```

```

In [12]: # lets check the features
terms = tfidf_vectorizer.get_feature_names()
terms[23361]

```

```

Out[12]: 'seed'

```

```
In [13]: # lets see the centroids into which clusters they belongs
         for i in range(5):
             print("\nCluster %d:" % i)
             for ind in order_centroids[i, :10]:
                 print('%s' % terms[ind])
```

Cluster 0:

yukos
china
virus
oil
microsoft
russian
spyware
gazprom
security
windows

Cluster 1:

said
growth
market
year
economy
sales
bank
company
prices
economic

Cluster 2:

mr
said
labour
people
election
blair
government
party
brown
minister

Cluster 3:

game

england
said
win
cup
club
match
team
players
play

Cluster 4:
film
best
awards
award
music
band
year
said
films
album

lets do the prediction and check the results

4. Visualization


- Lets plot the features in a 2D space.
- The dimension of features got from TfidfVectorizer is large (> 10,000)so we will do the dimensionality reduction using Principal Component Analysis.
- Here, PCA will transform the high dimensional features into 2 dimensions.

```
In [14]: def plot_clusters(reduced_features, reduced_cluster_centers, cls):  
    plt.scatter(reduced_features[:,0], reduced_features[:,1], c=cls.predict(features))  
    plt.scatter(reduced_cluster_centers[:, 0], reduced_cluster_centers[:,1], marker='x', s=150, c='b')  
  
    # reduce the features to 2D  
    pca = PCA(n_components=2, random_state=random_state)  
    reduced_features = pca.fit_transform(features.toarray())  
  
    # reduce the cluster centers to 2D  
    reduced_cluster_centers_mk = pca.transform(mkmeans.cluster_centers_)  
    reduced_cluster_centers_k = pca.transform(kmeans.cluster_centers_)
```


lets visualize using scatter plot.

- In the plot:
 - X = first dimension
 - Y = second dimension
- Different colors represent the cluster members.

```
In [15]: # MK Means  
  
plot_clusters(reduced_features, reduced_cluster_centers_mk, mkmeans )
```




```
In [16]: # K Means  
plot_clusters(reduced_features, reduced_cluster_centers_k, kmeans)
```



5. Evaluation

- Let us check the predictions.

```
In [17]: print("Prediction using KMeans")  
X = tfidf_vectorizer.transform(["The new social network is popular in cell phones"])  
predicted = kmeans.predict(X)  
print(predicted)
```

Prediction using KMeans
[2]

- The prediction label is [2] i.e. cluster 2. It is related with technology and also the examined text is also technology related. This prediction is correct.
- Further, we have to test the classifier with other text as well.

```
In [18]: print("Prediction using Mini Batch KMeans")  
X = tfidf_vectorizer.transform(["The new social network is popular in cell phones"])  
predicted = mkmeans.predict(X)  
print(predicted)
```

```
Prediction using Mini Batch KMeans  
[1]
```

Is the prediction accurate??

In []:

Other formal methods for evaluation:

- Since we have a labelled dataset is easier than the dataset without labelled dataset.
- `homogeneity_score`: values ranges from 0 and 1 where 1 indicates ideally homogeneous labeling.

```
In [19]: from sklearn.metrics import homogeneity_score  
hs_mk = homogeneity_score(df.label, mkmeans.predict(features))  
hs_k = homogeneity_score(df.label, kmeans.predict(features))  
  
print("Homogenity Mini KMeans: %f"%hs_mk)  
print("Homogenity KMeans: %f"%hs_k)
```

```
Homogenity Mini KMeans: 0.573055  
Homogenity KMeans: 0.619430
```

- In absence of labelled dataset, Cluster evaluation can be done using other metrics like Silhouette Coefficient (SC).
- SC uses mean intra-cluster distance (a) and the mean nearest-cluster distance (b) for each sample.
- Here, b is the distance between a sample and the nearest cluster that the sample is not a part of.
- The SC for a sample is $(b - a) / \max(a, b)$.
 - best value = 1
 - worst value = -1
 - near 0 values means overlapping clusters.
 - Negative values means bad clustering.

```
In [20]: from sklearn.metrics import silhouette_score  
ss_mk = silhouette_score(features, labels=mkmeans.predict(features))  
ss_k = silhouette_score(features, labels=kmeans.predict(features))  
  
print("silhouette Mini KMeans: %f"%ss_mk)  
print("silhouette KMeans: %f"%ss_k)
```

```
silhouette Mini KMeans: 0.009695  
silhouette KMeans: 0.009264
```

References:

- <https://sanjayasubedi.com.np/nlp/nlp-with-python-document-clustering/>
- <https://www.kaggle.com/jbencina/clustering-documents-with-tfidf-and-kmeans>
- <http://kavita-ganesan.com/extracting-keywords-from-text-tfidf/>
- <https://towardsdatascience.com/applying-machine-learning-to-classify-an-unsupervised-text-document-e7bb6265f52>

In []: