

Text Clustering Example ¶

- This example shows text Clustering as an example of unsupervised machine learning task.
- The main steps are:
 - Explore dataset
 - Data Preparation
 - Feature Engineering
 - Model Training
 - Performance Assessment
- Clustering: grouping similar objects together based on their inherent attributes.
- Being unsupervised approach there is no need of prelabelled datasets.
- Examples: KMeans, PAM, DBSCAN, Spectral clustering, etc.
- The choice of the algorithm mainly depends on the problem domain and the algorithm property. For example, to use KMeans we should have pre-knowledge of the number of clusters to obtain. However, if we do not know the number of clusters already, then DBSCAN is better.

Problem Definition

- As discussed above, the problem in context is unsupervised learning problem.
- We will use BCC news datasets available from <http://mlg.ucd.ie/files/datasets/bbc-fulltext.zip> (<http://mlg.ucd.ie/files/datasets/bbc-fulltext.zip>). A total of 2225 documents with five different news categories(i.e. business, entertainment, politics, sport, and tech) are available in the dataset.

```
In [1]: # Import necessary libraries

import numpy as np
import pandas as pd
from sklearn.cluster import MiniBatchKMeans
from sklearn.cluster import KMeans
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
```

1. Load and explore data

- The dataset contains 5 different categories in 5 different folders.

```
In [2]: from sklearn.datasets import load_files

# for reproducibility
random_state = 22

DATA_DIR = "./data/bbc/"
data = load_files(DATA_DIR, encoding="utf-8", decode_error="replace", random_state=random_state)
df = pd.DataFrame(list(zip(data['data'], data['target'])), columns=['text', 'label'])
```

It is necessary to do an exploratory data analysis in order to gain some insights from the data.

How is the data like?

```
In [3]: df.head()
```

Out[3]:

how many 'class label' present in the dataset??

```
In [4]: df['label'].unique()
```


Out[4]:

Is the dataset balanced??

If the dataset is balanced then, the dataset contains an approximately equal portion of each class.

```
In [5]: df['label'].value_counts().plot(kind='bar')
```

Out[5]:



2. Feature Engineering

- we will use TFIDF approach.
- TfidfVectorizer will convert all the documents to a matrix of TF-IDF features.
- With Tfidfvectorizer we will calculate the word counts, IDF and TFIDF values in a go.

```
In [6]: tfidf_vectorizer = TfidfVectorizer(stop_words="english")
tfidf_vectorizer.fit(df.text.values)
features = tfidf_vectorizer.transform(df.text.values)
```

Lets check some values from the first document

```
In [7]: # get the first vector out (for the first document)
first_vector_tfidfvectorizer=features[0]

# place tf-idf values in a pandas data frame
feature_df = pd.DataFrame(first_vector_tfidfvectorizer.T.todense(), index=features.get_feature_names(), columns=[0])
feature_df.sort_values(by=["tfidf"], ascending=False).head()
```

Out[7]:

Type *Markdown* and LaTeX: α^2

3. Model Building

- We will check a version of KMeans algorithm, i.e. Mini-batch KMeans (MB KMeans).
- MB KMeans is similar to the standard KMeans. But in the MB KMeans the computationally expensive step is performed on a random sample only instead of the entire data.
- We will set the number of clusters = 5 because our dataset contains categories.
- But, if we do not have knowledge of the different labels then we may have to using domain knowledge and other ways to get the optimal numbers of clusters.

```
In [8]: mkmeans = MiniBatchKMeans(n_clusters=5, random_state=random_state)
mkmeans.fit(features)
```

Out[8]:

lets check the predictions using labels_ attribute of the model.

```
In [9]: mkmeans.predict(features)
print mkmeans.labels_
print
print mkmeans.cluster_centers_
```

```
[4 2 4 ... 4 4 1]
```

```
[[1.34473036e-04 1.09906331e-02 0.00000000e+00 ... 0.00000000e+00
  0.00000000e+00 0.00000000e+00]
 [1.68041680e-04 1.05754629e-02 0.00000000e+00 ... 0.00000000e+00
  0.00000000e+00 0.00000000e+00]
 [2.47895846e-04 2.12690116e-03 0.00000000e+00 ... 0.00000000e+00
  4.60818743e-04 0.00000000e+00]
 [0.00000000e+00 5.00195763e-03 0.00000000e+00 ... 0.00000000e+00
  0.00000000e+00 0.00000000e+00]
 [4.67908112e-06 1.15377847e-02 4.45013653e-05 ... 4.09596975e-05
  0.00000000e+00 4.63992050e-04]]
```

Check KMeans

```
In [10]: kmeans = KMeans(n_clusters=5, init='k-means++', max_iter=100, n_init=1)
kmeans.fit(features)
kmeans.predict(features)
print kmeans.labels_
print
print kmeans.cluster_centers_
```

```
[1 3 1 ... 1 1 4]
```

```
[[0.00000000e+00 1.49656595e-02 0.00000000e+00 ... 0.00000000e+00
  0.00000000e+00 0.00000000e+00]
 [1.71467723e-04 1.21722706e-02 7.47173745e-05 ... 1.37541895e-04
  0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 9.52220007e-03 0.00000000e+00 ... 0.00000000e+00
  0.00000000e+00 0.00000000e+00]
 [1.59264506e-04 2.40800061e-03 0.00000000e+00 ... 0.00000000e+00
  4.79581592e-04 3.92171424e-04]
 [1.45232699e-04 1.02439695e-02 0.00000000e+00 ... 0.00000000e+00
  0.00000000e+00 0.00000000e+00]]
```

```
In [11]: # lets order the centroids

order_centroids = kmeans.cluster_centers_.argsort()[:, :-1]
order_centroids
```

Out[11]:

```
In [12]: # lets check the features
terms = tfidf_vectorizer.get_feature_names()
terms[23361]
```

Out[12]:

```
In [13]: # lets see the centroids into which clusters they belongs
for i in range(5):
    print("\nCluster %d:" % i)
    for ind in order_centroids[i, :10]:
        print('%s' % terms[ind])
```

Cluster 0:

growth
economy
sales
prices
said
economic
year
2004
dollar
bank

Cluster 1:

said
mr
government
new
company
year
film
uk
people
000

Cluster 2:

mr
labour
election
blair
party
said
brown
howard
government
tax

Cluster 3:

best
game
england
film
win
said
year
won
cup
play

Cluster 4:

mobile
people
said
technology
users
music
digital
software
phone
games

lets do the prediction and check the results

4. Visualization

- Lets plot the features in a 2D space.
- The dimension of features got from TfidfVectorizer is large (> 10,000)so we will do the dimensionality reduction using Principal Component Analysis.
- Here, PCA will transform the high dimensional features into 2 dimensions.


```
In [14]: def plot_clusters(reduced_features, reduced_cluster_centers, cls):  
    plt.scatter(reduced_features[:,0], reduced_features[:,1], c=cls.predictions)  
    plt.scatter(reduced_cluster_centers[:, 0], reduced_cluster_centers[:, 1], c=cls.predictions)  
  
    # reduce the features to 2D  
    pca = PCA(n_components=2, random_state=random_state)  
    reduced_features = pca.fit_transform(features.toarray())  
  
    # reduce the cluster centers to 2D  
    reduced_cluster_centers_mk = pca.transform(mkmeans.cluster_centers_)  
    reduced_cluster_centers_k = pca.transform(kmeans.cluster_centers_)
```

lets visualize using scatter plot.

- In the plot:
 - X = first dimension
 - Y = second dimension
- Different colors represent the cluster members.


In [15]: *# MK Means*

```
plot_clusters(reduced_features, reduced_cluster_centers_mk, mkmeans )
```



In [16]: *# K Means*

```
plot_clusters(reduced_features, reduced_cluster_centers_k, kmeans)
```



5. Evaluation

- Let us check the predictions.

```
In [17]: print("Prediction using KMeans")
X = tfidf_vectorizer.transform(["The new social network is popular in c
predicted = kmeans.predict(X)
print(predicted)
```

Prediction using KMeans
[4]

- The prediction label is [2] i.e. cluster 2. It is related with technology and also the examined text is also technology related. This prediction is correct.
- Further, we have to test the classifier with other text as well.

```
In [18]: print("Prediction using Mini Batch KMeans")
X = tfidf_vectorizer.transform(["The new social network is popular in c
predicted = mkmeans.predict(X)
print(predicted)
```

Prediction using Mini Batch KMeans
[1]

Is the prediction accurate??

In []:

Other formal methods for evaluation:

- Since we have a labelled dataset is easier than the dataset without labelled dataset.
- homogeneity_score: values ranges from 0 and 1 where 1 indicates ideally homogeneous labeling.

```
In [19]: from sklearn.metrics import homogeneity_score
hs_mk = homogeneity_score(df.label, mkmeans.predict(features))
hs_k = homogeneity_score(df.label, kmeans.predict(features))

print("Homogeneity Mini KMeans: %f"%hs_mk)
print("Homogeneity KMeans: %f"%hs_k)
```

Homogeneity Mini KMeans: 0.573055
Homogeneity KMeans: 0.547866

- In absence of labelled dataset, Cluster evaluation can be done using other metrics like Silhouette Coefficient (SC).
- SC uses mean intra-cluster distance (a) and the mean nearest-cluster distance (b) for each sample.
- Here, b is the distance between a sample and the nearest cluster that the sample is not a part of.

- The SC for a sample is $(b - a) / \max(a, b)$.
 - best value = 1
 - worst value = -1
 - near 0 values means overlapping clusters.
 - Negative values means bad clustering.

```
In [20]: from sklearn.metrics import silhouette_score
ss_mk = silhouette_score(features, labels=mkmeans.predict(features))
ss_k = silhouette_score(features, labels=kmeans.predict(features))

print("silhouette Mini KMeans: %f"%ss_mk)
print("silhouette KMeans: %f"%ss_k)
```

```
silhouette Mini KMeans: 0.009695
silhouette KMeans: 0.009480
```

References:

- <https://sanjayasubedi.com.np/nlp/nlp-with-python-document-clustering/> (<https://sanjayasubedi.com.np/nlp/nlp-with-python-document-clustering/>)
- <https://www.kaggle.com/jbencina/clustering-documents-with-tfidf-and-kmeans> (<https://www.kaggle.com/jbencina/clustering-documents-with-tfidf-and-kmeans>)
- <http://kavita-ganesan.com/extracting-keywords-from-text-tfidf/> (<http://kavita-ganesan.com/extracting-keywords-from-text-tfidf/>)
- <https://towardsdatascience.com/applying-machine-learning-to-classify-an-unsupervised-text-document-e7bb6265f52> (<https://towardsdatascience.com/applying-machine-learning-to-classify-an-unsupervised-text-document-e7bb6265f52>)

```
In [ ]:
```